

From Image Collections to Point Clouds with Self-supervised Shape and Pose Networks

K L Navaneet¹ Ansu Mathew¹ Shashank Kashyap¹ Wei-Chih Hung²
Varun Jampani³ R. Venkatesh Babu¹

¹Indian Institute of Science ²University of California, Merced ³Google Research

Abstract

Reconstructing 3D models from 2D images is one of the fundamental problems in computer vision. In this work, we propose a deep learning technique for 3D object reconstruction from a single image. Contrary to recent works that either use 3D supervision or multi-view supervision, we use only single view images with no pose information during training as well. This makes our approach more practical requiring only an image collection of an object category and the corresponding silhouettes. We learn both 3D point cloud reconstruction and pose estimation networks in a self-supervised manner, making use of differentiable point cloud renderer to train with 2D supervision. A key novelty of the proposed technique is to impose 3D geometric reasoning into predicted 3D point clouds by rotating them with randomly sampled poses and then enforcing cycle consistency on both 3D reconstructions and poses. In addition, using single-view supervision allows us to do test-time optimization on a given test image. Experiments on the synthetic ShapeNet and real-world Pix3D datasets demonstrate that our approach, despite using less supervision, can achieve competitive performance compared to pose-supervised and multi-view supervised approaches.

1. Introduction

3D object reconstruction is a long standing problem in the field of computer vision. With the success of deep learning based approaches, the task of single image based 3D object reconstruction has received significant attention in the recent years. The problem has several applications such as view synthesis and grasping and manipulation of objects.

Early works [4, 2, 3] on single image based 3D reconstruction utilize full 3D supervision in the form of 3D voxels, meshes or point clouds. However, such approaches require large amounts of 3D data for training, which is hard




Figure 1. **Single-image 3D Reconstructions.** Input image and corresponding projection from reconstructed 3D point clouds. We reconstruct the 3D output from a single input image using a completely self-supervised approach.

and expensive to obtain. Several recent works [23, 21] have focused on utilizing multi-view 2D supervision in the form of color images and object silhouettes as an effective alternative training protocol. A key component in these techniques is the differentiable rendering module that enables the use of 2D observations as supervision using reprojection consistency based losses. However, most of these approaches require multiple 2D view of the same 3D model along with the associated camera pose information in the training stage. This is a major limitation in applying these techniques in a practical setting where such supervisory data is difficult to obtain.

In this work, we set out to tackle a more challenging problem of learning 3D object reconstructions from image and corresponding silhouette collections. Given a collection of images and corresponding object silhouettes belonging to the same object category such as car, with just a single view from each object instance and no ground truth camera pose information, our goal is to learn 3D object reconstructions (Fig. 1). The proposed approach is practically useful and enables us to make effective use of the large amounts of 2D training data for learning 3D reconstructions. Since it is possible to easily obtain object silhouettes in the absence of ground truth masks, here we make the reasonable assumption that the image collection contains corresponding silhouettes. A key challenge in our training setting is to simultaneously learn both camera pose estimation and 3D

reconstruction while avoiding degenerate solutions. For instance, a degenerate solution for 3D reconstruction would be to just lift 2D pixels in a given image onto a 3D plane. Although such a flat 3D reconstruction perfectly explains a given image, that is obviously not a desired 3D shape. In this work, we introduce loss functions that are tailored towards simultaneous learning of the pose and reconstruction networks while avoiding such degenerate solutions. Specifically, we propose to use geometric and pose cycle consistency losses. To enforce geometric cycle consistency, we make use of the fact that multiple 2D views from the same 3D model must all result in the same 3D model upon reconstruction. However, note that these multiple 2D views are intermediate representations obtained in our framework utilizing just a single image per model. To correctly regress the pose values, we obtain projections from random viewpoints to enforce consistency in pose predictions. Motivated by the observation that the reconstruction performance improves remarkably when multiple 2D views are used for supervision, we aim to utilize additional images as supervision. However, since our problem setting limits the number of views from each 3D model to one, we effectively retrieve images from the training set with similar 3D geometry in a self-supervised manner. We utilize them as additional supervision in the form of cross-silhouette consistency to aid the training of pose and reconstruction networks.

Since our approach is self-supervised, we can adapt our network to obtain better reconstructions on a given test input image by performing additional optimization during inference. We propose regularization losses to avoid over-fitting on a single test sample. This ensures that the 3D reconstructions are more accurate from input viewpoint while maintaining their 3D structure in the occluded regions.

We benchmark our approach on the synthetic ShapeNet [1] dataset and observe that it achieves comparable performance to the state-of-the-art multi-view supervised approaches [16, 7]. We also evaluate our approach on the real-world Pix3D [18] dataset and show comparable or improved performance over a pose supervised approach [16]. We also present possible applications of our approach for dense point correspondence and 3D semantic part-segmentation. To the best of our knowledge, this is the first completely self-supervised approach for 3D point cloud reconstruction from image and silhouette collections.

To summarize, we make the following contributions in this work:

- We propose a framework to achieve single image 3D point cloud reconstruction in a completely self-supervised manner.
- We introduce cycle consistency losses on both pose and 3D reconstructions to aid the training of the pose and reconstruction networks respectively.

- We effectively mine images from geometrically similar models to enforce cross-silhouette consistency, leading to significantly improved reconstructions
- We perform thorough evaluations to demonstrate the efficacy of each component of the proposed approach on the ShapeNet dataset. We also achieve competitive performance to pose and multi-view supervised approaches on both ShapeNet and real-world Pix3D datasets.

2. Related Works

Single Image Based 3D Reconstruction Several learning based works in the recent past have tackled the problem of single image based 3D object reconstruction. The initial works [4, 2, 3, 19, 5, 12] make use of full 3D supervision in terms of ground-truth voxels or point clouds. Choy *et al.* [2] utilize multiple inputs for improved voxel reconstructions. Fan *et al.* [3] is one of the first works to learn point cloud reconstructions from images using a deeply learned network. They made use of set distance based losses to directly regress the 3D locations of the points. Mandikal *et al.* [13] extend [3] to predict point clouds with part segmentations using a part-aware distance metric calculation.

2D Supervised Approaches While the above works obtain promising results, they require ground truth 3D models as supervision which is complex and expensive to obtain. To overcome this, several works [23, 21, 22, 24, 11, 15, 6, 20, 7, 16, 9] have explored 2D supervised approaches utilizing 2D images, silhouettes, depth maps and surface normal maps. These works aim to develop ways to go from the 3D representation to the 2D projections in a differentiable manner in order to effectively back-propagate the gradients from the 2D loss functions to the reconstruction network. Yan *et al.* [23] achieve this on voxel based 3D representations by performing grid sampling of voxels to obtain foreground mask projections. Re-projection losses are used from multiple viewpoints to train the network. Similarly, Tulsiani *et al.* [21] use a differentiable ray consistency based loss to reconstruct not only the shape information, but also features like color. The work is extended in [20] where a multiple-view consistency based loss is formulated to simultaneously predict 3D camera pose and object reconstructions. Motivated by the computational and performance advantages offered by point clouds, a number of works have sought to design rendering modules for projecting 3D points. Insafutdinov and Dosovitskiy [7] and Navaneet *et al.* [15, 16] develop differentiable projection modules to project points and corresponding features on to the 2D plane, enabling training on 2D representations like silhouettes, depth maps, images and part segmentations.




Figure 2. Approach Overview. We propose a cycle consistency based approach to obtain 3D reconstructions from a collection of images and their corresponding foreground masks. An encoder-decoder architecture based network is used to regress the 3D coordinates of the point cloud reconstruction \hat{P} . A pose network is used to obtain 3D camera pose predictions $\hat{\mathbf{v}}$ from the input image. DIFFER [16] is used to render the reconstruction in the predicted viewpoint. Additionally, reconstructions are also projected from randomly sampled poses to obtain k projections which are again used to reconstruct k point clouds \hat{P}^k . We enforce a 3D cycle consistency loss on \hat{P} and \hat{P}^k to train \mathcal{N}_{rec} . Similarly the randomly sampled poses and corresponding projections are considered as pseudo ground truth labels to enforce pose cycle consistency loss. The red dashed arrows in the diagram indicate the proposed losses.

Weakly Supervised Approaches Among the weakly supervised approaches, [8, 14, 10, 16, 20, 7] are the closest to ours. Mees *et al.* [14] utilize mean 3D object models to learn 3D reconstructions in a self-supervised manner. Li *et al.* [10] generate 3D models using a self-supervised approach, but do not perform reconstruction from RGB images. In SSL-Net [17], 3D models are used to pre-train one of the networks before performing self-supervised reconstruction. To the best of our knowledge, we are the first to obtain colored 3D point cloud reconstructions from just a collection of images and corresponding silhouettes.

3. Approach

We aim to obtain 3D point cloud reconstruction from a single image in a self-supervised setting. To this end, we propose a learning based approach with an encoder-decoder architecture based network to predict the reconstructions. Let I be the image input to the network, M the foreground object mask and $\hat{P} \in \mathbb{R}^{N \times 3}$ the corresponding point cloud reconstruction obtained using the reconstruction network \mathcal{N}_{rec} (refer Fig. 2). N is the number of points in the reconstructed point cloud. In the absence of ground truth 3D models, all our supervisory data, which is the set of input images and corresponding silhouettes, lies in the 2D domain. In order to utilize these 2D observations to train the network, we would need to project the reconstructed point cloud on to the 2D image plane. We use the differentiable projection modules proposed in DIFFER [16] and CAPNet [15] to obtain color and mask projections respectively from a given viewpoint. The viewpoint \mathbf{v} associated with the input image is characterized by azimuth and elevation values of the camera in the 3D space placed at a

fixed distance from the object. We use another encoder network $\mathcal{N}_{\text{pose}}$ to obtain the viewpoint prediction $\hat{\mathbf{v}}$. The reconstructed point cloud is projected from the predicted viewpoint using the differentiable projection module to obtain 2D image and mask predictions \hat{I} and \hat{M} respectively. If the predicted viewpoint and reconstructions are correct, the 2D projections will match the input image. To enforce this, we use the losses proposed in DIFFER [16] to optimize both the reconstruction and pose prediction networks. Specifically, we use the following image (\mathcal{L}_I) and mask (\mathcal{L}_M) loss functions:

$$\mathcal{L}_I = \frac{1}{hw} \sum_{i,j} \|I_{i,j} - \hat{I}_{i,j}\|_2^2 \quad (1)$$

$$\mathcal{L}_{\text{bce}} = \frac{1}{hw} \sum_{i,j} -M_{i,j} \log \hat{M}_{i,j} - (1 - M_{i,j}) \log (1 - \hat{M}_{i,j}) \quad (2)$$

$$\begin{aligned} \mathcal{L}_{\text{aff}} &= \sum_{i,j} \min_{(k,l) \in M_+} ((i-k)^2 + (j-l)^2) \hat{M}_{i,j} M_{k,l} \\ &\quad + \sum_{i,j} \min_{(k,l) \in \hat{M}_+} ((i-k)^2 + (j-l)^2) M_{i,j} \hat{M}_{k,l} \end{aligned} \quad (3)$$

$$\mathcal{L}_M = \mathcal{L}_{\text{bce}} + \mathcal{L}_{\text{aff}} \quad (4)$$

where h, w are the height and width of the 2D observations respectively. M_+ and \hat{M}_+ are sets of pixel coordinates of the ground truth and predicted projections whose values are non-zero. In this formulation, the predictions by the reconstruction and pose networks rely heavily on each other. Since the predicted viewpoint is used in projection, the reconstruction network can predict correct 3D models that consistently match the input image only if the

pose predictions are accurate. Similarly, since the pose network parameters are optimized using projection losses, the predicted pose values will be correct only if the reconstructions are reasonable. In such a situation, the network can collapse to degenerate solutions. For instance, the predicted viewpoint can be constant regardless of the input and the 3D reconstruction can be planar. The networks would still achieve zero loss as long as they reproduce the input image from the predicted constant viewpoint. To avoid such degenerate solutions, we propose novel cycle consistency losses to train both reconstruction and pose networks.

3.1. Geometric Cycle Consistency Loss

We propose geometric cycle consistency loss to train the reconstruction network (Fig. 2) to avoid degenerate reconstructions. The reconstructed point cloud \hat{P} is projected from k randomly sampled viewpoints $\{\mathbf{v}^i\}_1^k$. Let $\{\hat{I}^i\}_1^k$ be the corresponding image projections. These images are used as input to the reconstruction network \mathcal{N}_{rec} and the corresponding reconstructed point clouds $\{\hat{P}^i\}_1^k$ are obtained. Since each of the projections and the input image are all associated with the same 3D object, the corresponding point clouds must all be consistent with each other. To enforce this, we define the geometric cycle consistency loss as follows:

$$\mathcal{L}_G = \sum_{i=1}^k d_{\text{Ch}}(\hat{P}, \hat{P}^i) \quad (5)$$

where $d_{\text{Ch}}(\cdot, \cdot)$ denotes the Chamfer distance between two point clouds. The reconstruction network is trained using a combination of the mask and image losses and the geometric cycle consistency loss.

$$\mathcal{L}_{\text{rec}}^{\text{total}} = \alpha(\mathcal{L}_I + \mathcal{L}_M) + \beta\mathcal{L}_G \quad (6)$$

3.2. Pose Cycle Consistency Loss

The projection based losses form weak supervisory signals to train the pose prediction network. While there is no direct pose information available for the input images, the projected images and corresponding pose pairs $\{\hat{I}^i, \mathbf{v}^i\}_1^k$ can be considered as pseudo ground-truth pairs for training the pose network. We input the image projections $\{\hat{I}^i\}_1^k$ to the pose prediction network $\mathcal{N}_{\text{pose}}$ to obtain the corresponding pose predictions $\{\hat{\mathbf{v}}^i\}_1^k$ (Fig 2). The corresponding viewpoints $\{\mathbf{v}^i\}_1^k$ are then used as ground-truth to train $\mathcal{N}_{\text{pose}}$. The pose loss is obtained as follows:

$$\mathcal{L}_{\text{pose}} = \frac{1}{k} \sum_{i=1}^k |\mathbf{v}^i - \hat{\mathbf{v}}^i| \quad (7)$$

The final training objective for the pose network is a combination of pose cycle consistency loss and image and mask losses (Eq. 1 and 4). This ensures that the pose loss is



Figure 3. **Sample k-nearest neighbours.** We utilize our single-view trained reconstruction network to obtain k-nearest neighbour samples from the train set. Note that the neighbours have different poses and have different color distribution, but have similar 3D shape which provides us with additional information on the geometry of the object.

dependent on the pose predictions of the input image, while simultaneously being optimized with a stronger supervision using the projected images.

$$\mathcal{L}_{\text{pose}}^{\text{total}} = \gamma(\mathcal{L}_I + \mathcal{L}_M) + \rho\mathcal{L}_{\text{pose}} \quad (8)$$

3.3. Nearest Neighbours Consistency Loss

Earlier works [15] demonstrate that even just a single additional view as supervision during training significantly improves the reconstruction quality. However, as mentioned previously, assuming the presence of such multi-view images during training curtails the practical utility and prevents the applicability on real-world single image datasets. In order to remain in the constrained setting, but improve reconstructions with the use of multiple image supervision, we propose mining images from the training set which belong to similar 3D models. For every input image, we find the closest neighbours such that they have similar underlying 3D shapes, and use projection consistency based loss, termed ‘nearest neighbours consistency loss’, to assist the training of the network. To find the nearest neighbours in the 3D domain in a self-supervised fashion, we need features which embed the 3D shape information. Utilizing features from networks trained on 2D tasks (for e.g., classification on ImageNet dataset), would provide neighbours which are similar in color and viewpoint, but not necessarily in 3D shape. Alternatively, to quantify the 3D similarity, we consider the encoded features of our proposed reconstruction network. Nearest neighbours from training set are obtained by comparing the Euclidean distances in the encoded feature space. Sample nearest neighbour images are shown in Fig. 3. We observe that the retrievals are similar in shape and have diversity in terms of pose and color. During training, nearest neighbours of the input image are utilized as additional supervision. The neighbour images are passed through $\mathcal{N}_{\text{pose}}$ to obtain the corresponding poses. The reconstructed point cloud obtained from the input image is projected from these viewpoints. We then enforce silhouette loss in Eq. 4 on these projections using the

ground-truth silhouettes of the neighbour images. This is possible since the geometry of the input and the neighbours are similar and thus the projections from the input model closely match those from the neighbours. Note that the loss is enforced using only masks and not the color images since the neighbours might have different color distribution. The mask losses are summed over n neighbours to get the total nearest neighbours loss. This is used in addition to the losses mentioned in Eq. 1 and 4 to train the reconstruction network.

$$\mathcal{L}_{\text{NN}} = \sum_{i=1}^n \mathcal{L}_M^i \quad (9)$$

3.4. Symmetry Loss

Since all the object categories we consider in our experiments have a minimum of one plane of symmetry, we further regularize the network to obtain symmetric reconstructions with respect to a pre-defined plane. Without loss of generality, let us assume that the point clouds are symmetric with respect to the xz-plane. Then, the symmetry loss is given by:

$$\mathcal{L}_{\text{sym}} = d_{\text{Ch}}(\hat{P}^+, \hat{P}^-) \quad (10)$$

where \hat{P}^+ is the set of points in \hat{P} with positive y values and \hat{P}^- is the reflection about the xz-plane of the points in \hat{P} with negative y values. The symmetry loss helps in obtaining improved geometry of reconstructions consistent with the ground truth and avoids overfitting to the input image. Due to the absence of ground truth pose values, the co-ordinate system for the predicted camera poses is not pre-determined. The choice of plane of symmetry in enforcing symmetry loss can also help align the reconstructions to a predefined canonical pose. The total reconstruction loss with nearest neighbours and symmetry losses is as follows:

$$\mathcal{L}_{\text{rec}}^{\text{total}} = \alpha(\mathcal{L}_I + \mathcal{L}_M) + \beta\mathcal{L}_G + \eta\mathcal{L}_{\text{NN}} + \kappa\mathcal{L}_{\text{sym}} \quad (11)$$

3.5. Inference Stage Optimization (ISO)

Our self-supervised approach, which relies only on the input images and corresponding object silhouettes for training, is ideally poised for instance specific optimization during inference. At inference, we predict both the 3D point locations and the input image viewpoint. To obtain highly corresponding reconstructions, we aim to minimize the difference between the input and the projected image (from predicted viewpoint) during inference. To ensure that the reconstructions are not degraded in the regions occluded in the input image, we employ additional regularization. Note that while CAPNet [15] too performs inference stage optimization, unlike our work, the authors assume known viewpoint. The regularization loss formulation is as follows:

$$\mathcal{L}_{\text{reg}} = d_{\text{ch}}(\hat{P}, \hat{P}_O) \quad (12)$$

where \hat{P} and \hat{P}_O are the initial and optimized point clouds. We also use the symmetry loss as an additional form of regularization to enable the network to optimize for the regions in the point cloud visible in the input image while suitably modifying the points corresponding to the occluded regions. The objective function during ISO is given by:

$$\mathcal{L}_{\text{ISO}} = \alpha(\mathcal{L}_I + \mathcal{L}_M) + \lambda(\mathcal{L}_{\text{reg}}) + \kappa(\mathcal{L}_{\text{sym}}) \quad (13)$$

4. Experiments

4.1. Implementation Details

We use a two-branch network to simultaneously obtain shape and color reconstructions. Separate models are used for training on each object category. The number of projections, k is set to four and the number of points in reconstructed point cloud to 1024. Adam optimizer with a learning rate of 0.00005 is used for training the network. The hyperparameters α , β , γ and ρ are set to 100, 10^4 , 1 and 1 respectively. Architecture details, additional details on hyperparameter settings and training schedules are provided in the supplementary material. We publicly release the code.¹

4.2. Datasets

ShapeNet [1]: ShapeNet is a curated set of synthetic 3D mesh models. We sample points on the surface of the meshes to obtain the corresponding point clouds for evaluation. To create the set of input images, we render the mesh models from a single random view per object instance. All the experiments are performed on the representative car, chair and airplane (denoted as aero) categories.

Pix3D [18]: Pix3D is a repository of aligned real-world image and 3D model pairs. The dataset exhibits great diversity in terms of object shapes and backgrounds and is highly challenging. We consider the chair category of Pix3D in our experiments. Since the dataset is small, we only perform evaluation on the Pix3D dataset.

We use the train/val/test splits provided by DIFFER [16] in all our experiments. For ease of comparison, all the Chamfer and EMD metrics are scaled by 100.

4.3. Evaluation Methodology

Since point clouds are unordered representations, we use Chamfer distance and earth mover's distance (EMD) to evaluate the reconstructions. For evaluation, we randomly sample 1024 points from the reconstructions if they contain higher number of points. The Chamfer distance between two point clouds P and \hat{P} is defined as $d_{\text{Chamfer}}(P, \hat{P}) = \sum_{x \in P} \min_{y \in \hat{P}} \|x - y\|_2^2 + \sum_{x \in \hat{P}} \min_{y \in P} \|x - y\|_2^2$. EMD between two point clouds is defined as $d_{\text{EMD}}(P, \hat{P}) = \min_{\phi: P \rightarrow \hat{P}} \sum_{\alpha \in P} \|\alpha - \phi(\alpha)\|_2$

¹Code is available at https://github.com/val-iisc/ssl_3d_recon

where $\phi(\cdot)$ is a bijection from P to \hat{P} . For the pose unsupervised approaches, the models are aligned using a global rotation matrix obtained by minimizing the Chamfer error on the validation set. To evaluate color metrics, we project each reconstruction from 10 randomly sampled viewpoints and compute the L_2 distance using the ground-truth images. We report the median angular error and accuracy in the pose prediction evaluation. In addition, the pose metrics are also calculated by utilizing the ground truth orientation. The predicted point cloud is ‘flipped’ (rotated by 180°) if the error is more than 90° .

4.4. Baseline Approaches

We compare the proposed approach with two state-of-the-art approaches on 2D supervised single image based 3D point cloud reconstruction. Specifically, we use the following variants of the works:

DIFFER: DIFFER [16] proposed a differentiable module to project point cloud features on to the 2D plane, which enables it to utilize input images for training. Note that DIFFER utilizes ground truth pose values for the input image and hence has a higher degree of supervision compared to our approach. Codes and settings provided by the authors are used to train the network.

ULSP: Insafutdinov *et al.* [7] proposed a multi-view consistency based unsupervised approach for point cloud reconstruction. While the approach does not make use of ground truth pose values, it requires multiple images and their corresponding foreground masks from different viewpoints per 3D object instance. Hence, the work is not directly comparable to our approach which uses just a single image per model. To remain as close as possible to this setting, we train ULSP with supervision from two views per model using the code provided by the authors.

ULSP_Sup: We consider a variant of ULSP [7] with ground truth camera pose supervision. Similar to DIFFER, this is trained with one input viewpoint per 3D model.

We also provide comparison with two variants of the proposed approach - ‘Ours-CC’ and ‘Ours-NN’. Ours-CC is trained only with the cycle consistency losses while NN consistency loss is used in addition in Ours-NN.

4.5. Effect of Cycle Consistency Losses

We first analyze the role of the proposed consistency losses in improving the reconstructions in a self-supervised setting (Table 1). In the absence of both \mathcal{L}_G and $\mathcal{L}_{\text{pose}}$ (Ours-No-CC), the network fails to learn meaningful 3D reconstructions. When both the cyclic losses are employed (Ours-CC), we observe that the network learns the underlying 3D shapes of the objects and thus results in effective reconstructions. We present detailed ablations for individual loss components in the supplementary material.

Method	Chamfer			EMD		
	Car	Chair	Aero	Car	Chair	Aero
Ours-No-CC	10.33	21.84	15.06	18.32	23.40	16.12
Ours-CC	6.39	13.58	8.66	6.42	16.46	12.53

Table 1. **Effect of Consistency Loss.** We evaluate the effect of the proposed consistency losses on reconstruction metrics. The network fails to train in the absence of the consistency losses in the self-supervised setting.

Method	Chamfer			EMD		
	Car	Chair	Aero	Car	Chair	Aero
DIFFER	6.35	9.78	5.67	6.03	16.21	9.9
DIFFER + \mathcal{L}_G	5.63	9.23	5.58	5.35	13.07	9.44
ULSP_Sup	6.64	10.49	5.70	6.89	10.93	7.43
ULSP_Sup + \mathcal{L}_G	6.13	10.0	7.37	5.83	10.24	9.99

Table 2. **Portability of Geometric Consistency.** Using our geometric consistency loss atop supervised approaches results in significant gains in reconstruction performance.

We also demonstrate the utility of the proposed geometric losses in the pose supervised setting for single image based 3D reconstructions. Specifically, we use the proposed loss \mathcal{L}_G atop pose supervised DIFFER and ULSP_Sup to optimize the corresponding reconstruction networks. Table 2 suggests that the geometric loss can significantly improve the performance of existing supervised approaches as well.

4.6. Reconstruction Results

Quantitative and qualitative comparisons of the proposed self-supervised approach with other multi-view and pose supervised approaches on the ShapeNet dataset are provided in Table 3 and Fig. 4 respectively. The performance of our approach is comparable to those utilizing higher levels of supervision. For the baseline approaches, we observe that pose supervised ULSP_Sup is marginally better than the two-views supervised ULSP in the case of chairs and airplanes and significantly better in the case of cars. Our car reconstruction metric is close to the supervised ULSP network and is better than other approaches. Notably, while we use the same projection module and projection consistency losses as in DIFFER, we outperform the pose supervised DIFFER in most of the quantitative metrics. This demonstrates the utility of the additional cycle and nearest neighbour consistency loss for reconstruction and pose prediction. The addition of nearest neighbour significantly boosts the reconstruction performance, particularly in the case of the more challenging chair category. In the car and airplane categories, there is apparent visual improvement in the shape and spread of points with the use of nearest neighbours. While we are able to effectively capture the geometry of the object, points are sparsely distributed in the




Figure 4. Comparisons on ShapeNet. We provide comparison with both pose and multi-view supervised approaches on ShapeNet. Our approach is on par with the supervised approaches in terms of correspondence of the reconstruction to the input image. Our car reconstructions have significantly better shape and uniformity in points compared to the supervised approaches.

thin regions such as legs in the case of chairs. However, we can observe similar sparse point distributions in the case of DIFFER [16]. We also present qualitative (Fig. 5) and quantitative (in supplementary) results on inference stage optimization. The reconstructions have greater correspondence with the input image as observed in the silhouettes before and after optimization in Fig. 5. Reconstruction metrics indicate that the point clouds are preserved in regions not observed in the test input. Additional qualitative results, ablations on symmetry and nearest neighbours consistency loss and failure cases are provided in the supplementary.

To show the adaptability of our approach to real-world datasets, we evaluate it on the Pix3D dataset. Note that since the dataset consists of very few models, we perform evaluation of the networks trained on ShapeNet dataset. For synthetic to real domain adaptation, we train on ShapeNet dataset with the input images overlaid with random natural scene backgrounds. Our approach performs comparably to the pose supervised DIFFER approach both quantitatively (Table 4) and qualitatively (Fig. 6).

Fig. 7 presents qualitative results on color prediction on ShapeNet dataset. For effective evaluation, we project each ground truth and predicted model from 10 randomly sampled viewpoints and calculate the channel-wise \mathcal{L}_2 loss between them. Our reconstructions result in greater visual correspondence with the input image, particularly in the case of cars. Quantitative results are provided in the supplementary.

4.7. Pose Prediction Results

Table 5 presents median error and accuracy of our pose prediction network. We report results both with ('flip') and without ('No-flip') the use of ground-truth orientation. Ours-CC achieves high accuracy on the car category. How-




Figure 5. Inference stage optimization (ISO). Optimization during inference results in greater correspondence to the input image. Regularization is employed to maintain the shape in regions occluded in the input image.




Figure 6. Comparisons on Pix3D. Since both DIFFER and the proposed approach are trained on ShapeNet and evaluated on Pix3D, the correspondence to input in reconstructions is lower compared to that on ShapeNet. However, our reconstructions have marginally better shape and point spread compared to the supervised DIFFER approach.

ever, in the chair category, where there exists higher ambiguity, Ours-CC performs significantly worse. Due to the existence of multiple planes of symmetry in certain airplane models, the network often predicts the incorrect orientation, as observed in the high median error. But when the ground truth orientation is used to calculate the metrics, such conflicts are resolved leading to significantly better metrics. In all the categories, we observe that the pose metrics reliably

Method	Pose	Views	Chamfer			EMD		
			Car	Chair	Aero	Mean	Car	Chair
ULSP_Sup	Yes	1 view	5.4	9.72	5.91	7.01	4.78	10.18
DIFFER	Yes	1 view	6.35	9.78	5.67	7.27	6.03	16.21
ULSP	No	2 views	7.02	9.87	5.96	7.62	7.99	10.56
Ours-CC	No	1 view	6.39	13.58	8.66	9.54	6.42	16.46
Ours-NN	No	1 view	5.48	10.91	7.11	7.83	4.95	14.93
							12.53	11.8
							11.07	10.31

Table 3. **Reconstruction Metrics on ShapeNet.** Despite being self-supervised, lacking the input pose values and with just the input image as supervision, we perform comparably to or even outperform other state-of-the-art approaches requiring higher degree of supervision.

Method	Chamfer	EMD
DIFFER	14.33	16.09
Ours-NN	14.52	15.82

Table 4. **Reconstruction Results on Pix3D.** We evaluate both the pose supervised DIFFER and our approach on the real-world Pix3D [18] dataset. Our self-supervised approach performs comparably to the pose supervised one and adapts well to the real-world dataset.




Figure 7. **2D Color Projections.** Our colored projections have greater visual correspondence with the input images compared to the supervised DIFFER approach.

Categ.	Method	Median Error		Accuracy	
		No-flip	Flip	No-flip	Flip
Car	Ours-CC	7.58	5.54	74.07	94.4
	Ours-NN	6.85	5.55	75.87	93.4
Chair	Ours-CC	41.86	33.78	41.45	45.72
	Ours-NN	19.69	17.79	59.14	64.16
Aero	Ours-CC	88.29	38.53	20.99	40.74
	Ours-NN	43.36	19.52	42.34	60.74

Table 5. **Pose Metrics on ShapeNet.** Pose metrics are remarkably good for the car category for both our approaches. In the challenging chair and airplane categories, use of nearest neighbours (Ours-NN) significantly boosts the predictions.

improve upon the introduction of nearest neighbor consistency (\mathcal{L}_{NN}), further highlighting the need for such a loss. We also observe that the pose and reconstruction metrics are correlated and thus incorrect prediction in either of them significantly affects the other.

4.8. Point Correspondences and Part Transfer

In our reconstructions, we observe that points with similar indices in the regressed points have spatial correspon-




Figure 8. **Part Transfer.** Semantic part segmented ground truth and reconstructed point clouds. Correspondences between the reconstructed point clouds are used for consistent part segmentation transfer across models.

dence even though we do not explicitly enforce it. We use a colored UV map to visualize the point correspondences (see Supplementary for more details). We utilize this correspondence for the task of single-shot semantic part segmentation. We use a single ground-truth part-segmented model to transfer part labels across all models based on point indices. Results (Fig. 8) indicate that our network is effective in obtaining 3D part segmentations using just a single ground-truth model.

5. Conclusion

We propose a self-supervised approach for single image based 3D point cloud reconstruction. We develop novel geometric and pose cycle consistency losses to effectively train our reconstruction and pose networks in a self-supervised manner. Through the use of training images with similar 3D shape, we mimic the effect of training with multi-view supervision using a single-view dataset. We benchmark our reconstruction, color and pose prediction networks on the ShapeNet dataset, achieving comparable performance to pose and multi-view supervised approaches. The role of all the proposed losses is thoroughly analyzed. We further demonstrate the utility of our approach through reconstruction results on the real-world Pix3D dataset and qualitative results on possible applications like dense point correspondence and 3D part segmentation. In future, we like to address the issue of sparse point predictions in thin structures and further improve the reconstruction quality.

Acknowledgement This work was supported by Sudha Murthy Chair Project, Pratiksha Trust, IISc.

References

- [1] Angel X Chang, Thomas Funkhouser, Leonidas Guibas, Pat Hanrahan, Qixing Huang, Zimo Li, Silvio Savarese, Manolis Savva, Shuran Song, Hao Su, et al. Shapenet: An information-rich 3D model repository. *arXiv preprint arXiv:1512.03012*, 2015. 2, 5
- [2] Christopher B Choy, Danfei Xu, JunYoung Gwak, Kevin Chen, and Silvio Savarese. 3D-r2n2: A unified approach for single and multi-view 3D object reconstruction. In *European Conference on Computer Vision*, pages 628–644. Springer, 2016. 1, 2
- [3] Haoqiang Fan, Hao Su, and Leonidas Guibas. A point set generation network for 3D object reconstruction from a single image. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, volume 38, 2017. 1, 2
- [4] Rohit Girdhar, David F Fouhey, Mikel Rodriguez, and Abhinav Gupta. Learning a predictable and generative vector representation for objects. In *European Conference on Computer Vision*, pages 484–499. Springer, 2016. 1, 2
- [5] Thibault Groueix, Matthew Fisher, Vladimir G. Kim, Bryan Russell, and Mathieu Aubry. AtlasNet: A Papier-Mâché Approach to Learning 3D Surface Generation. In *CVPR*, 2018. 2
- [6] Paul Henderson and Vittorio Ferrari. Learning to generate and reconstruct 3d meshes with only 2d supervision. *arXiv preprint arXiv:1807.09259*, 2018. 2
- [7] Eldar Insafutdinov and Alexey Dosovitskiy. Unsupervised learning of shape and pose with differentiable point clouds. In *NeurIPS*, 2018. 2, 3, 6
- [8] Angjoo Kanazawa, Shubham Tulsiani, Alexei A Efros, and Jitendra Malik. Learning category-specific mesh reconstruction from image collections. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 371–386. 3
- [9] Hiroharu Kato, Yoshitaka Ushiku, and Tatsuya Harada. Neural 3d mesh renderer. In *CVPR*, 2018. 2
- [10] Xiao Li, Yue Dong, Pieter Peers, and Xin Tong. Synthesizing 3d shapes from silhouette image collections using multi-projection generative adversarial networks. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2019. 3
- [11] Chen-Hsuan Lin, Chen Kong, and Simon Lucey. Learning efficient point cloud generation for dense 3D object reconstruction. In *AAAI Conference on Artificial Intelligence (AAAI)*, 2018. 2
- [12] Priyanka Mandikal, K L Navaneet, Mayank Agarwal, and R Venkatesh Babu. 3D-LMNet: Latent embedding matching for accurate and diverse 3d point cloud reconstruction from a single image. In *Proceedings of the British Machine Vision Conference (BMVC)*, 2018. 2
- [13] Priyanka Mandikal, K L Navaneet, and R Venkatesh Babu. 3D-PSRNet: Part segmented 3d point cloud reconstruction from a single image. In *3D Reconstruction Meets Semantics Workshop (ECCVW)*, 2018. 2
- [14] Oier Mees, Maxim Tatarchenko, Thomas Brox, and Wolfram Burgard. Self-supervised 3d shape and viewpoint estimation from single images for robotics. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2019. 3
- [15] K L Navaneet, Priyanka Mandikal, Mayank Agarwal, and R Venkatesh Babu. CAPNet: Continuous approximation projection for 3d point cloud reconstruction using 2d supervision. In *AAAI*, 2019. 2, 3, 4, 5
- [16] K L Navaneet, Priyanka Mandikal, Varun Jampani, and R Venkatesh Babu. DIFFER: Moving beyond 3d reconstruction with differentiable feature rendering. In *CVPR Workshops*, 2019. 2, 3, 5, 6, 7
- [17] Ran Sun, Yongbin Gao, Zhijun Fang, Anjie Wang, and Cengsi Zhong. Ssl-net: Point-cloud generation network with self-supervised learning. *IEEE Access*, 7:82206–82217, 2019. 3
- [18] Xingyuan Sun, Jiajun Wu, Xiuming Zhang, Zhoutong Zhang, Chengkai Zhang, Tianfan Xue, Joshua B Tenenbaum, and William T Freeman. Pix3d: Dataset and methods for single-image 3d shape modeling. In *CVPR*, 2018. 2, 5, 8
- [19] Maxim Tatarchenko, Alexey Dosovitskiy, and Thomas Brox. Octree generating networks: Efficient convolutional architectures for high-resolution 3D outputs. In *CVPR*, 2017. 2
- [20] Shubham Tulsiani, Alexei A Efros, and Jitendra Malik. Multi-view consistency as supervisory signal for learning shape and pose prediction. In *CVPR*, pages 2897–2905, 2018. 2, 3
- [21] Shubham Tulsiani, Tinghui Zhou, Alexei A Efros, and Jitendra Malik. Multi-view supervision for single-view reconstruction via differentiable ray consistency. In *CVPR*, 2017. 1, 2
- [22] Jiajun Wu, Yifan Wang, Tianfan Xue, Xingyuan Sun, Bill Freeman, and Josh Tenenbaum. Marrnet: 3D shape reconstruction via 2.5 d sketches. In *Advances In Neural Information Processing Systems*, pages 540–550, 2017. 2
- [23] Xincheng Yan, Jimei Yang, Ersin Yumer, Yijie Guo, and Honglak Lee. Perspective transformer nets: Learning single-view 3D object reconstruction without 3D supervision. In *Advances in Neural Information Processing Systems*, 2016. 1, 2
- [24] Rui Zhu, Hamed Kiani Galoogahi, Chaoyang Wang, and Simon Lucey. Rethinking reprojection: Closing the loop for pose-aware shape reconstruction from a single image. In *Computer Vision (ICCV), 2017 IEEE International Conference on*, pages 57–65. IEEE, 2017. 2